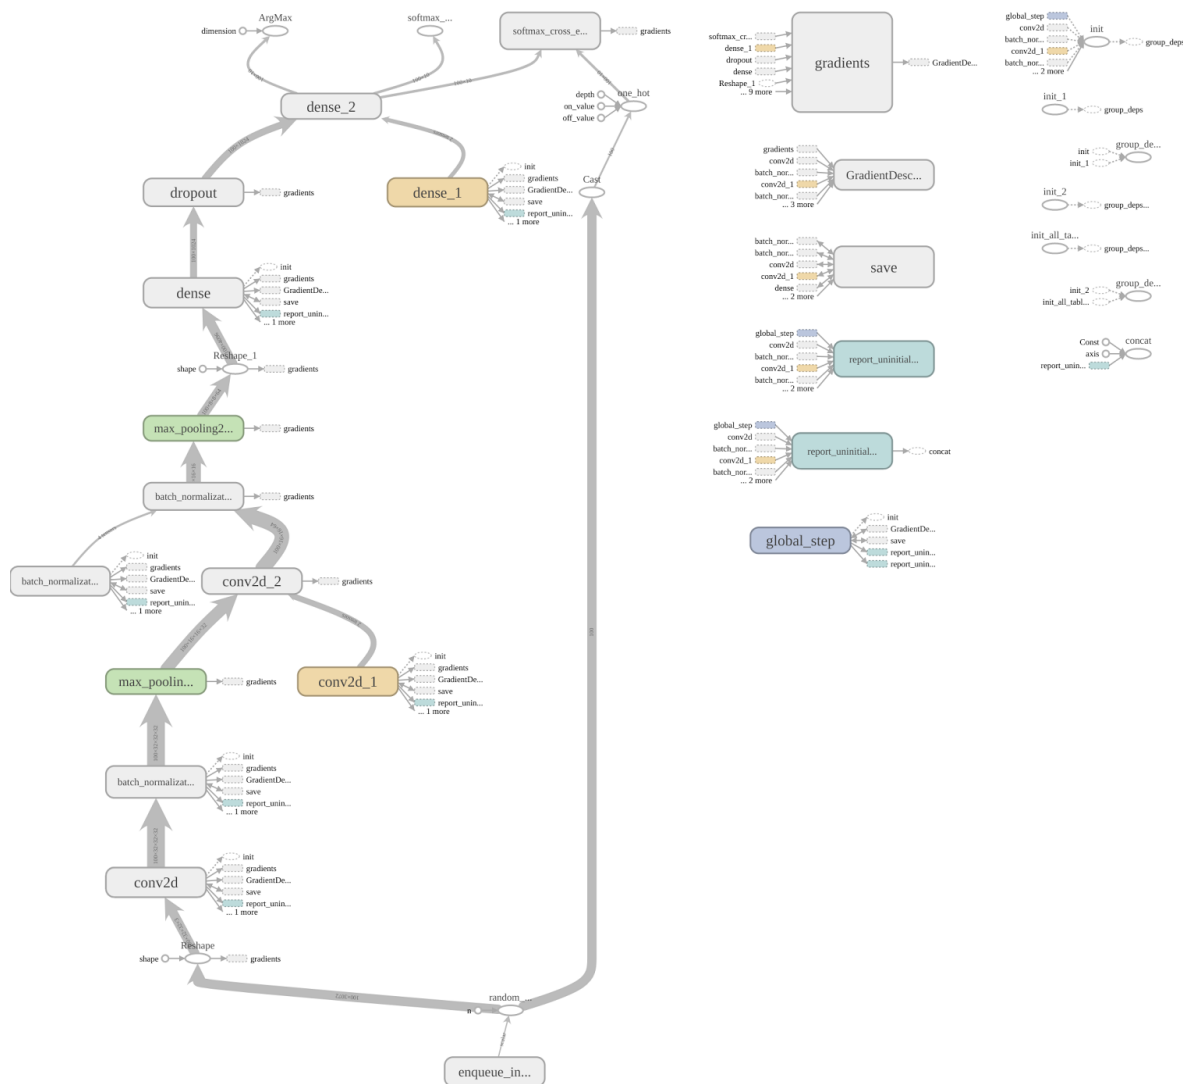


Homework 2

QUESTION 1:

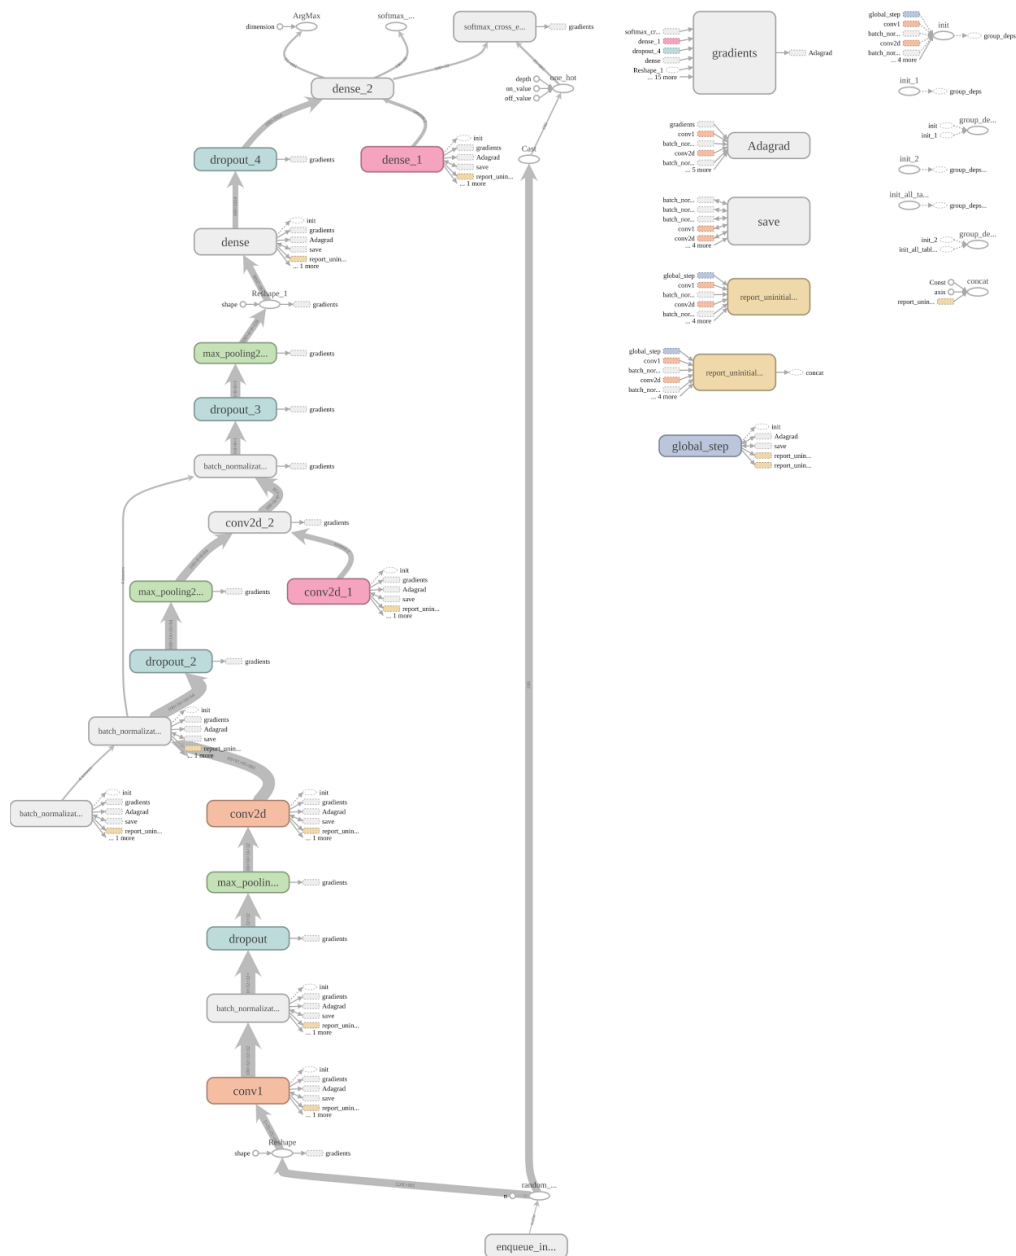
Architecture:

The architecture consists of two convolutional layers with 32 and 64 filters respectively. Each convolutional layer is followed by batch normalization and max pooling. This is then followed by a dense layer. We perform dropout after the dense layer with a dropout rate of 0.4.



Part 1 Validation Accuracy: **72.74%**

Architecture:



Steps:

1. Different Optimizers

- ADAGrad Optimizer
- ADAM Optimizer
- RMSPROP Optimizer

Best result:

ADAGradOptimizer

2. Dropouts

- We tried different configurations of

Best result:

Dropouts on all layers

LAYERS	DROPOUTS
Layer 1	0.2
Layer 2	0.3
Layer 3	0.4
Layer 4	0.3

3. Different Layers

- 5 Layer Network with [64,64,512,32] filters each layer respectively with a dense layer (1024 units)
- 5 Layer Network with [64,64,32,32] filters each layer respectively with a dense layer(1024 units)
- 4 Layer Network with [32,64,128] filters each layer respectively with a dense layer(1024 units)
- 4 Layer Network with [32,64,64] filters each layer respectively with a dense layer(1024 units)

Best Result:

4 Layer Network with [32,64,128] filters each layer respectively with a dense layer(1024 units)

While trying these experiments, our validation accuracy varied between 72%-78.08%.

Final Architecture:

- Training size: We did not use data augmentation. We used the original data which is 45K (5K validation data)
- Train Time: 16 minutes on a GPU
- 4-layer network: [32,64,128] filters with 3 convolutional layers, with a final dense layer with 1024 units.
- Each of the convolutional layers were followed by batch normalization, max pooling and dropout (dropout rates: 0.2, 0.3, 0.4 respectively).


Validation accuracy improvement with each optimization:

- Adagrad Optimizer (vs Gradient Descent Optimizer in part 1): **~2%**
- Change in layer architecture (from 2 Convolutional layers [32,64] to 3 Convolutional layers [32, 64, 128]): **~2%**
- Adding dropout after each layer with dropout rates 0.2, 0.3 and 0.4: **~2%**


Part 2 - Best Training Accuracy: **99.17%**

Part 2 - Best Validation Accuracy: **78.44%**

Convolution Layer 1




Convolutional Layer 2



Convolutional Layer 3

conv_layer3



Last Fully Connected Layer

dense_layer1

